Test the promise. Verify the result.

Solvia quality control process and current release evidence

OpenAI's US-developed GPT-OSS 120B (openai/gpt-oss-120b), served through Together AI. The model drafts editable facts; deterministic application services decide prices, ownership, payment and permit issuance.

Verification layer	Verified result	What the result establishes
Final backend	603/603 passed 76.503 seconds	Final-source C#/xUnit: ownership, draft retention, shared vehicles, prices, quote locking, provider contracts, receipts and PDFs.
Main baseline	132 distinct covered	TypeScript Playwright; 44 smoke executions are a subset. Full run plus targeted rerun; original timeout retained on page 3.
Tablet baseline	24 distinct covered	Portrait/landscape Chromium/WebKit. Original timeouts and any targeted resolution remain explicit. Emulation is not physical-device or actual Safari certification.
Live AI baseline	12/12 passed 168.487 seconds	Separate opt-in vehicle/schedule cases. Accepted Together output requires pinned model identity, schema and grounding. No payment is started.
Actual Sandbox payments	Historical 20-check matrix	Four campuses; 14 completed captures / $945 gross test funds, two free terms and 16 checked PDFs. Two additional provider negatives and two SDK/no-JS smokes remain separate.

Broad browser/model/guide baseline: 9de5bd40f72a43565a56b5c6b6ffdceba08a41db. Final backend and 24/24 focused UI checks cover the later shared-vehicle guard. The broad browser/model suites and guide were not all rerun on that delta.

Baseline browser coverage is 156 distinct cases: 132 main plus 24 tablet; live AI and final focused UI remain separate. Reruns add no distinct case. Simulated callbacks/order creation do not prove capture. Historical payments remain historical.

Three subagents and the coordinating agent reviewed the reported workflows and all 21 scanned pages. These are AI-assisted technical reviews, not independent certification, human usability research, adoption or measured ROI.

Invariants before features

Risk	Release acceptance rule
Lost or overwritten answers	Optional schedule text can be blank. Explicit revise and browser Back retain the current draft. Unchanged vehicle text cannot overwrite reviewed corrections. A changed sentence receives fresh extraction.
Stale or foreign draft	Drafts are bound to the server-resolved demo person and flow. Switching people cannot copy a stale form into another person. Protected browser drafts expire; saved applications remain available separately.
Unreviewed schedule	Changing dates, weekdays, weeks or campus invalidates the old comparison. The server rechecks comparison freshness even without JavaScript. Continuing with daily permits starts only one planned day.
Wrong fee or changed quote	Only published catalog cents and the server category set the fee. Missing prices block payment. An existing PayPal order freezes its purchase snapshot; changed unbound fees require renewed review before creating an order.
Stale-tab purchase	CheckoutReviewSnapshot binds owner/application, campus/category, permit type/name, term/dates, cents and actual vehicle fields. Start reloads authoritative facts within its transaction and refuses a changed fingerprint before a new order. This is freshness, not authorization. Payment status/order/capture are excluded for unchanged retries.
Owner-safe 429 recovery	Rate rejection retains the current server-resolved DraftOwner, owned FlowId and entered answers. A stale form cannot restore another person. Retry-After precedes model invocation; normal owner/flow checks still apply.
False payment or duplicate purchase	Paid issuance requires COMPLETED capture, matching amount/currency and a capture ID. Retry references remain durable. Existing order-bound drafts cannot be revised into a different purchase.
Invented model success	Pinned response.model, strict final-content JSON, grounding, response-size limits and elapsed deadlines precede the Together label. Invalid/unavailable responses use a labelled fallback. No weights are fine-tuned.
Confusing payment controls	Official SDK options appear directly on Check your total. The SDK decides method eligibility. Merely loading the page creates no order. Protected no-JavaScript/failed-SDK continuation remains available.
Misleading permit proof	Paid and Active are separate. PDFs and receipts preserve their own plate, dates and payment identifiers. A separate historical receipt is never presented as payment for a fresh unpaid draft.

The model has no payment tools. US developer origin does not establish US-only processing, retention settings or zero data retention. Credentials remain server-side; prose and raw provider transport errors stay out of public evidence.

Reproduce. Repair. Review.

1. Reproduce with fictional input and source; separate app/provider/harness faults. 2. Add a regression and valid boundaries. 3. Run focused and broader checks; keep real-provider evidence separate and reconcile purchases before retry. 4. Inspect UI/PDFs, reconcile reviewers, retain failures and freeze/verify the authorized source.

Retained failures make the record useful

Backend candidate: 551/582 passed, 31 failed. Twenty-seven EF required-navigation exceptions and four downstream HTTP 500-versus-409 results exposed the new snapshot reload path. The fix loads/reloads Vehicle and PermitType by authoritative foreign keys after refreshing the application; it avoids IsLoaded resets and tracker clearing. A warm-FK-change regression was added.

Next backend run: 582/583 passed. An old test expected a later concurrency exception; corrected reload instead rejected a $40 quote against the $30 catalog before provider access. The expectation now requires zero provider calls, unchanged 4,000 cents and no order. Baseline then passed 583/583; final-delta full backend is separately shown on page 1. Failed TRXs remain retained.

Main baseline: original 131/132 passed; targeted 1/1 passed on the same application source. Final UI assertions passed before the 30-second overall budget expired ahead of cleanup completion; the trace proves no specific hang. Only that test's budget became 60 seconds; assertion waits stayed 10 seconds.

Tablet baseline: original 22/24 passed; targeted 2/2 resolved the same failed identities. Two 60-second timeouts had unfinished assertions and multi-hour trace-event gaps. Durations are retained but unreliable; host interruption is not proven. Production was unchanged within this baseline rerun.

The first final-guard UI run passed 22/24: two garage assertions used reload after an invalid POST, replaying the rejected form. The corrected test opens garage/Edit with GET to inspect saved state. No production change was needed; the same 24 cases were rerun cleanly. The failed report remains retained.

Earlier retained findings include planner POST-history/Back failure (repaired with protected GET review), numeric plate/state recognition initially 6/9 (then 124/124 extraction checks), and restricted test-host IPC attempts that ran no tests. These are not relabelled as clean passes.

Practical limits and next release gates

This fictional Sandbox demo uses researched CSV fees and demo dates, not live campus eligibility or authorization. Interactive buyer-wallet approval, verified delivered webhooks, refunds, live money and physical devices remain outside the verified scope.

Public YouTube publication, clean full-path judge access and participant eligibility/rights answers remain gates. No final submission, external certification or award is claimed.

Reproduction and current source: https://github.com/jlhpacheco/solvia

Evidence index: docs/TESTING.md; docs/QUALITY-CONTROL.md; docs/review/USER-REVIEW-2026-10-20.md; docs/review/Solvia-repair-verification.json.